

Project Report - Entregas Caracol Lda

1. Solution Description

The problem consists of calculating the number of distinct paths between all pairs of intersections (A, B) in a one-way map (Directed Acyclic Graph – DAG) and assigning truck routes based on this value.

The solution, implemented in C++, is fully iterative and avoids exponential/sub-exponential approaches. The steps are:

1. Topological Sort (Kahn's Algorithm)

- Linearizes the graph to ensure that when visiting a node u , all predecessors have already been processed.
- Complexity: $O(N + E)$.

2. Batched Dynamic Programming

- Source nodes are processed in blocks of up to 128 vertices (BATCH_SIZE).
- For each batch, a serialized array `batch_counts[u * BATCH_SIZE + k]` accumulates the number of paths between source (`start_base + k`) and node u .
- A `batch_reachable` vector and a token mechanism avoid clearing the entire matrix between batches; only the used slices are re-initialized (Lazy Reset).
- Propagation follows the topological order, summing counts modulo M and marking destinations as reachable.

3. Truck Assignment and Output

- For each reachable pair (S, T) , we calculate $truck(S, T) = 1 + (\text{paths}(S, T) \bmod M)$.
- Routes for trucks in $[m_1, m_2]$ are stored in vectors of pairs, later sorted lexicographically.

High-Level Pseudocode

```
KahnTopoOrder(G):
    compute in_degree
    queue q = vertices with in_degree 0
    topo = []
    while q not empty:
        u = pop(q); topo.append(u)
        for v in adj[u]:
            if --in_degree[v] == 0: push(q, v)
    return topo

Solve(G):
    topo = KahnTopoOrder(G)
    for each batch B of sources in topo order:
```

```

init_batch_slices(B)
for u in topo starting at first element of B:
    if u not touched by B: continue
    record deliveries for reachable pairs (source in B, u)
    propagate counts from u to each v in adj[u]
sort and print routes per truck

```

2. Theoretical Analysis

Time Complexity

Topological sorting is $O(N + E)$. Batched propagation executes, for each source, a traversal over the reachable subgraph: on average it is $O(N + E)$ per batch, resulting in a global complexity of

$$O(N \times (N + E))$$

In the dense extreme ($E \approx N^2$) this behaves as $O(N^3)$; in sparse graphs it approaches $O(N^2)$. Empirical tests confirm that the solution is suitable for expected limits and avoids exponential/sub-exponential growth.

Space Complexity

- **Graph (Adjacency List):** $O(N + E)$.
- **Auxiliary Structures (DP):** Vectors `batch_counts`, `batch_reachable`, `visit_token`, `topo_order` occupy $O(N \times \text{BATCH_SIZE})$; with `BATCH_SIZE = 128` this remains linear in practice.
- **Route Storage (Output):** In the worst case, all pairs (A, B) are valid. There are $N(N - 1)/2$ possible pairs. Thus, the space to store output is $O(N^2)$.

Total Space Complexity: $O(N^2 + E)$.

3. Experimental Evaluation

Experimental Results

Timings were measured on a MacBook Pro (macOS, Apple M-series CPU) running `./project` on DAGs generated by `python3 gen_test.py N` (standard density 0.3). Each instance is fed directly into the binary, and time is obtained with `time.perf_counter`. Full values are in `timings.csv`.

N	Time (s)	Ratio ($T_N/T_{N/2}$)
100	0.006	-
200	0.016	2.6x
400	0.055	3.4x

N	Time (s)	Ratio ($T_N/T_{N/2}$)
800	0.224	4.1x
1600	0.964	4.3x
3200	4.320	4.5x

Analysis: As N doubles, time grows between 3x–4.5x, reflecting the $O(N \times (N + E))$ term. Generated graphs become dense as N grows, so $E \propto N^2$ and behavior approaches $O(N^3)$, as predicted. Still, the solution easily processes instances up to $N = 3200$, comfortably meeting the time limit.

Automated tests include official examples, edge cases (minimal graph, disconnected components, multiple instances in same file, modular M scenarios), and larger random graphs, ensuring robustness and output correctness.